Actor Critic Methods: From Paper to Code

Monte Carlo Control Problem

Action Value Functions

- Without a model can’t transition to next valuable state
- Replace V with Q and use first visit MC
- How to handle explore exploit dilemma?

Exploring Starts

Pair random A, S → exploring starts

Pick random A at episode start

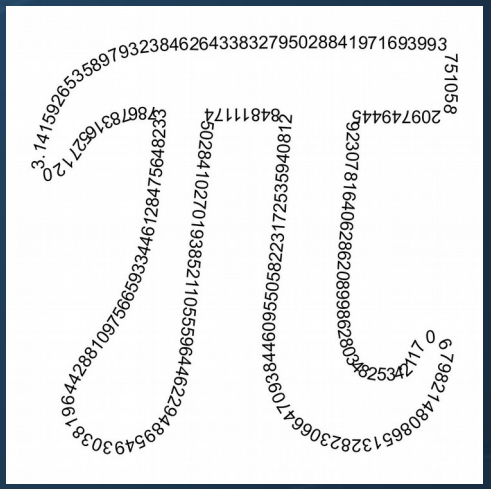

Gives good coverage of state and action space

Relaxing the Assumption of E.S.

- E.S. too limiting for some environments
- Use epsilon soft action selection

Exploratory action $\rightarrow \frac{\epsilon}{|A(s)|}$

Greedy action $\rightarrow 1-\epsilon+\frac{\epsilon}{|A(s)|}$

- Greed increases over time

Algorithm Overview

Initialize Q(s,a) arbitrarily for all s, a; terminal $\rightarrow$ 0

Initialize arbitrary epsilon soft policy

Initialize list of Returns(s,a) for all states and actions

Repeat for large number of episodes:

- Generate episode using policy
- For each state s and action a in the agent's memory:
  - Calculate the return that followed first visit to s, a
  - Append return G to list of Returns(s,a)
  - Update Q as the average of Returns(s,a)
- For each state s in the agent's memory:

$$A^* \leftarrow \underset{a}{argmax}\, Q(s,a)$$

$$\pi(a|s) = \begin{cases} 1-\epsilon+\epsilon/|A(s)| & if\ a = A^* \\ \epsilon/|A(s)| & if\ a \neq A^* \end{cases}$$

200,000 games; plot cum. win ratio over 1000 games $\epsilon \approx 0.001$

Conclusion

- Performance not too bad
- Able to do it without exploring starts


TACHYMETER
60
65
70
80
90
100
120
150
180
250
500
OMEGA
PROFESSIONAL
60
40
20
30
20
10
12
9
3
6